

Decide, build, verify and last with AI

A Practical Guide
for Entrepreneurs
and Managers



Sylvain Tremblay

Decide, Build, Verify, and Last with AI

A Practical Guide for Entrepreneurs
and Managers

Sylvain Tremblay

ST Publishing

Copyright © 2026 Sylvain Tremblay. All rights reserved.

This document is provided free of charge for informational purposes only. No part of this document may be reproduced or distributed in any form — electronic, mechanical, photocopy, or otherwise — without the written permission of the author or publisher, except for brief quotations in the context of articles or reviews.

The information contained in this document reflects the author's experience and research. While all efforts have been made to ensure accuracy, the author and publisher disclaim any responsibility for errors, omissions, or consequences resulting from the use of this information.

For any requests for permission or additional information, please contact: info@editio-nsst.com

Published in Quebec, Canada.


Table of content

The Gap is You.	1
AI isn't the problem. How you decide to use it is.	
1. What Actually Changed	3
The Cost of a First Draft Collapsed	
The Bottleneck Moved	
What This Actually Means	
2. The Decision Framework	6
How to Decide	
The Five Criteria	
What This Lets You Decide	
3. A Real Case	9
Example: Writing Product Descriptions	
4. What's Coming	11
The Two Books That Follow	
No Magic Tool	13

The Complete Series	15
Three Resources to Start Now	18
About the author	21

The Gap is You.

AI isn't the problem. How you decide to use it is.




*The Gap Triangle: decide what you keep, delegate the rest,
verify before it costs you.*

*The Gap Triangle: the core framework for deciding what to keep,
what to delegate to AI, and how to verify before it costs you.*

You have two problems with AI. The first is deciding what to hand off to it. The second—less visible but more expensive—is verifying that what it produces won't sabotage you.

Written by an entrepreneur who tested this on his own projects, with the real numbers, including the ones that don't favor him.

What you're about to read isn't an opinion. It's a method. A framework for deciding: what actually gets delegated to AI, and what absolutely cannot be.



1

What Actually Changed

The Cost of a First Draft Collapsed

Ten years ago, writing six hundred words with structure and definitions was a job. Someone spent days on it. You paid accordingly.

Today, that same work takes minutes.

This isn't a fifteen-percent improvement. It's a change in economic category.

Everything whose cost was driven by someone has to sit down and produce an initial version has collapsed:

- A sales proposal
- A product description
- A meeting summary

- A preliminary analysis
- A translation
- A contract template

You see the distinction? It's not that we write better. It's that you don't have to pay someone to write at all anymore. The first version appears from nothing in two minutes. It's not perfect, but it exists.

The Bottleneck Moved

For the entire history of office work, the bottleneck was production capacity. You hired to produce. You measured what got produced. Decisions and verification happened in the margins.

That's changed.

When something stops being rare, it stops being the bottleneck. The bottleneck shifts to the next constraint—which was already there, just invisible.

Now you can get five versions of a proposal in an hour. The real work is choosing which one. Knowing why you chose it. Having confidence that one of them won't shoot you in the foot without you seeing it.

And that's where it gets complicated. Because verifying what AI produces isn't about a quick glance anymore. It becomes an operation you have to design, document, build tools around, and protect.


What This Actually Means

We produce more. Much more. But we don't move faster.

The documents flow out. The queue gets longer. The bottleneck has simply shifted from one invisible place to another invisible place.

The real question isn't: Which tool should I choose?

The real question is: What do we actually keep? Because everything else can be handed off—but not without risk. And you carry that risk.



2

The Decision Framework

How to Decide

The question can we automate this? has an easy answer: yes, almost always. That's not the useful question.

The useful question is: What happens if it gets it wrong, and will I know?

This question has an objective answer. It depends on five criteria. Two of them have absolute veto power—if you fail on these two, the rest doesn't matter.

The Five Criteria

Cost of Failure. Not probability—real cost, actually counted. A misspelled name in a newsletter: you see it, fix it, it costs an hour

of work. A wrong amount on an invoice a customer receives: you lose the sale and the customer's trust. Two completely different universes.

Verifiability. Does an error show itself? How long does it take to spot it? There are three levels:

- A program catches it automatically (systematic verification)
- A human spots it in a quick read (needs a pair of eyes)
- Only someone who already knows the right answer can spot it (invisible otherwise)

Frequency. How many times per month. This criterion doesn't decide if it's possible—it decides how you do it. A task three times per year? Don't build an automated flow. Write clear instructions and save them. You'll read them again before each use. A task three times per week? Then yes, that deserves tooling.

Tacit Knowledge. How much of this work depends on things nobody has written down? To find out: ask the person who does this work how they decide. If they answer "it depends" more than twice in a row, you have your answer. High tacit knowledge.

Reversibility. If a wrong output slips through, can you fix it without consequences? An internal document can be changed. An email sent, a post published, a payment made, a contractual obligation signed: once it's out, it's out.


What This Lets You Decide

Here's the simple rule: if tacit knowledge is high, the answer is

document first, regardless of the other criteria. You can't delegate what you haven't written down.

For everything else—tacit knowledge low or moderate—these five criteria create a decision framework that holds.

The full chapter of this book shows you how to apply this to four real situations, with numbers backing it up. You'll see how the same criteria lead to different conclusions depending on the context.



3

A Real Case

Example: Writing Product Descriptions

You manage an e-commerce catalog with 300 products. You launch 20 new ones each month. Someone spends two hours per week just writing the descriptions.

The naive question: "We could just let AI write all that, right?"

The answer using the framework:

Cost of Failure: Two types of errors. A vague description that doesn't hook the customer = you lose the sale. Amount: dozens of dollars in revenue. A description that misses a crucial detail (material, size, exact color) = customer orders, receives the wrong product, returns it. Real cost: \$50 in shipping + return processing + relationship damage.

Verifiability: You read a description in thirty seconds and know if it sells. Immediate. You also have a metric: Is the return rate going up? That tells you.

Frequency: 20 products per month = 240 per year. Frequent enough to justify a system and tools.

Tacit Knowledge: Very high. Your typical customer on this product? Someone renovating their garden who wants a color that lasts 5 years. AI generates a beautiful, generic description that could work for ten different products. It doesn't know that the detail that *sells* is durability. It thinks it should talk about texture, color, price. It's wrong.


Reversibility: Complete. A bad description? Change it in two clicks. No permanent damage.

What This Gives You: You ask AI to generate 5 versions. You pick one. You add the detail that sells: "Lasts 5 years without maintenance. Unlike standard paint, which peels after 2-3 years." Total time: 30 minutes instead of 75.

Is this revolutionary? No. But it's real: you save 45 minutes per product $\times$ 20 products per month = 15 hours monthly. For an e-commerce operation, that's significant.

The Risk? Publishing the AI-generated descriptions as-is without adding the detail that sells. A generic description doesn't shock anyone—it just sells less. You lose the advantage you actually had: knowing your customer better than anyone.

This is a case where AI saves time. But only if you keep control over what makes you different.



4

What's Coming

The Two Books That Follow

You've just seen what it takes to decide.

The other two books in this series aren't repeating themselves. Each stands alone, but together they form a complete method.

AI in Action — Where It Really Works. Real Cases. Real Costs.

Tools change fast. Method doesn't.

Once you know what to build, you need to know how. Not the marketing around the tools—your hands in it. Which block to write, which tool to pick, which script to run, which safeguard to build in.

Nobody Decided This. — AI Ethics from the Owner's Desk.

But what's actually acceptable?

Four true stories from July-August 2026 where someone used AI

in a way that nobody had decided to allow. What happened. What it cost. How to prevent it.

No Magic Tool

What you just read is an assessment. You've seen how the criteria shift depending on context. How a good decision in one case isn't a good decision in another. How *you* have to make the call.

We could stop there. Say "here's the framework, figure it out."

But that's not enough.

Because a framework is like a compass. Useful. But if you don't know where you're going, a compass won't get you there.

So there are four questions this teaser hasn't answered yet.


Question 1: Where do I start? What can AI really do for my business, and where should I actually begin?

Question 2: Once I decide what to delegate to AI and what to keep, what do I actually do? Where does this fit in my company? What are the real use cases that'll make me money, not just waste my time?

Question 3: How do I build the verification? It's easy to say "you need to check." How do I actually do it? Which tool? Which process? Which script?

Question 4: What can still go wrong that I haven't seen coming? Because "AI misunderstood" isn't a problem. "AI misunderstood *and* my team didn't catch it *and* the customer received it"—that's a problem.

The four books that follow answer these four questions. Each one stands alone. None of them repeat each other. Together, they're a path forward.



The Complete Series

Book 1 — AI in Action

Where It Actually Works. Real Cases. Real Costs.

This book isn't "AI for Marketing." It's your marketing—where AI creates real impact and where it creates noise. Then your customer service, your data, your content creation—department by department.

The angle: you'll see on each case why it works or why it costs. And more importantly, what it *actually* costs—not in "hours saved" but in real money and real risk.

For: You who wants to see what works elsewhere before deciding what you'll do at home.

Available now in French and English.

Book 2 — The Gap is You.

Decide, Build, Verify, and Last with AI

You have the framework now. This book shows you how to deploy it. Not as a checklist. As a process.

How to ask the right questions. How to evaluate risks that the framework itself doesn't catch. Where to put the emergency brake. How to fight the perfectionist's paralysis ("we can never be 100% sure").

Every chapter = one real situation + how we decided + what it cost.

For: You who makes decisions. You who needs to convince a budget or a team. You who knows there's an answer, but you don't know yet how to find it.

Available now.

Book 3 — Hands in the AI

Write, Verify, Automate: The Hands-On Guide

Once you know what to build, you need to know how. Not the theory. Hands on it.

Which block to write so AI understands what you want. Which tool to pick among the three that would work. How to write the script that actually verifies instead of the script that "looks like it verifies." How to fix an instruction when it breaks at 3 AM.

This is a book to read with your hands on the keyboard. You'll copy things, adapt them, test them on your own cases.

For: You who already knows what to build and who's the technical person on your team. Or who's ready to be.

Coming soon.

Book 4 — Nobody Decided This.

AI Ethics from the Owner's Desk

We're using AI in ways that nobody had decided to allow. And it costs.

Four true stories from July-August 2026. In each one, someone used AI in a way that seemed right, then it ended badly. Not "AI hallucinated." It's: customer data exposed. Reputation damaged. Client filing a complaint. Or trust dead.

Not abstract ethical frameworks. Just: this happened, it cost this, here's how to prevent it.

For: You who says yes or no to AI. And who knows that yes can really go wrong.

Coming soon.

Start Here

stpublishing.ca


You'll find each book there. Access to the decision grid. And the supporting resources that follow.

Three Resources to Start Now

The Decision Grid — Printable Version

It's the one you saw here. In A4 format. Post it on your desk. Fill it out for the next thing someone wants to hand off to AI.

Decision Grid: Five Criteria



Use this decision grid to evaluate any task before delegating it to AI.

Download the Grid (PDF)

The Seven Most Common Mistakes — And How to See Them Coming

You read one case. We've seen these seven pitfalls in dozens of cases. They look like nothing. They cost a lot.

How to recognize them. How to prevent them at your place.

Seven Common Mistakes



[Print and post on your desk](#)
Click the image for the complete PDF

Seven common mistakes to avoid when delegating to AI — print and keep handy.

Download the Checklist (PDF)

Subscribe to Updates

When a book launches. When we add a case to the grid. When there's something important.

→ [Subscribe](#)

The Gap is You.

Decide, Build, Verify, and Last with AI

You've read enough to know if this is for you.

If it is, the rest is waiting.

ST PUBLISHING

Books for entrepreneurs and decision-makers who make real decisions.

stpublishing.ca

info@stpublishing.ca

© 2026 ST Publishing

About the author



Sylvain Tremblay earned his university degree in management information systems in 1992 and never left the field.

He started out in consulting at CGI, a large Canadian IT services firm. In 1996, he moved to Dolmen Logiciels, a small Quebec tax-software company, which was later absorbed into Intuit Canada — the Canadian division of the publisher of Quick-

en, QuickBooks, and TurboTax. The acquisition meant moving to Edmonton, Alberta, where Intuit Canada was based at the time. He spent 11 years there, from late 1996 to July 2007 — long enough for three of his four children to be born there. He came back with firsthand experience inside a Silicon Valley company, and with the conviction that ways of working are local and contingent: what's taken for granted in one place is a rule nowhere else.

He went out on his own when he got back. His consulting firm was born that same year, closely followed by Global-Lan, in web hosting and internet services, then by 55+ YOGA, co-founded in 2020, which offers adapted online yoga classes to people fifty and older and to those living with a physical limitation.

That last one doesn't sell technology — it depends on it entirely: a hosting platform, online course sales, billing, marketing, and various automations. That's exactly the profile of this book's reader — a small business whose trade lies elsewhere, and whose operations depend on tools that have to be chosen, connected, and maintained. He's actually used AI on his own processes before ever talking about it to anyone.

A consultant recommends; an owner pays. This book is written from that second seat.

Since generative AI arrived, he's been working on a single question: what actually works in a business of 10, 20, or 50 people. He tests, compares, builds complete chains, and runs them on his own projects before talking to anyone about them. The four situations that run through this book are his own, complete with their real numbers — including the ones that don't flatter him.

He trains and advises small business leaders and managers, and speaks on the subject. He published *AI in Action* in 2025, *The Gap is You.* in 2026 and is working on *Hands-On AI*, the technical companion to this same approach, as well as *Nobody Decided That.*, which asks the question of AI ethics seen from the owner's seat.

He lives in Quebec.

For permission requests or media inquiries: press@stpublishing.ca

For speaking engagements or consulting/training: consultation@stpublishing.ca

For bulk orders or other requests: contact@stpublishing.ca



Sylvain Tremblay is an IT consultant and entrepreneur with over thirty years of experience. He has founded and led several businesses, which has given him direct knowledge of small business constraints: tight budgets, lean teams, decisions with no safety net.

Since generative AI arrived, he tests rather than reads the promises. His books come from this hands-on practice: *AI in Action*, *The Gap is You*, *Hands in the AI*, and *Nobody Decided That*. This booklet is their introduction.

You have AI tools. The problem is the decision that surrounds them—not the tool itself.

This booklet isn't a complete book: it's a preview of a method that four books explore in depth. You'll find how to decide between what you delegate to AI and what you keep, how verification becomes a fortress instead of a weakness, how to build a process that actually holds. No products named, no screenshots—only the grids, the thresholds, and the decisions that nobody else puts into numbers.

This is the introduction. The four books that follow dig deeper.

For small business owners, managers, and consultants ready to move from "what can we do?" to "what do we decide to do?"

